



KubeCon



CloudNativeCon

Europe 2022

WELCOME TO VALENCIA





KubeCon



CloudNativeCon

Europe 2022

Debugging With Ephemeral Containers

Aaron Alpar, Kasten by Veeam



Ephemeral Containers

Debugging with Ephemeral Containers



Aaron Alpar
MTS, *Kasten by Veeam*




About this Presentation

- Introduction
- Ephemeral Containers
- Setup and Tools
- Debugging containers in K8s
- Linux Namespaces
- New Pod properties
- Creating Ephemeral Containers with custom `securityContext`
- Distroless



Kind



- Use Kind in this presentation (<https://kind.sigs.k8s.io>)
- Local Kubernetes environment
- Used for Kubernetes automated testing
- Uses Docker

Postgres

- Open source database
- Good target for examples



Setup: Kind

```
$ kind create cluster
```

```
Creating cluster "kind" ...
```

```
✓ Ensuring node image (kindest/node:v1.23.4)
```

```
✓ Preparing nodes
```

```
...
```

Not sure what to do next? 😓 Check out

<https://kind.sigs.k8s.io/docs/user/quick-start/>

```
$
```



Setup: Postgres



```
$ helm repo add bitnami https://charts.bitnami.com/bitnami
"bitnami" has been added to your repositories
```

```
$ kubectl create namespace pg
namespace/pg created
```

```
$ helm install pg bitnami/postgresql --namespace="pg"
```

```
NAME: pg
```

```
LAST DEPLOYED: Tue Apr 12 09:35:44 2022
```

```
...
```

```
PGPASSWORD="$POSTGRES_PASSWORD" psql --host 127.0.0.1 -U postgres -d
postgres -p 5432
```


Setup: Check versions

```
$ kubectl version --short
```

```
...
```

```
Client Version: v1.24.0
```

```
Kustomize Version: v4.5.4
```

```
Server Version: v1.23.4
```



Setup

```
$ kubectl get pods -n pg
```

NAME	READY	STATUS	RESTARTS	AGE
pg-postgresql-0	1/1	Running	0	30s



Ephemeral Containers

- Dynamically allocated containers
- Useful for debugging
- Beta in 1.23



Debug without Ephemeral Containers

- Ephemeral containers start with debugging



Debug without Ephemeral Containers

- Ephemeral containers start with debugging
- Debug by attaching (exec)

```
kubectl exec -it -n pg pg-postgresql-0 \  
-c postgresql -- /bin/sh
```



Debug without Ephemeral Containers

- Ephemeral containers start with debugging
- Debug by attaching (exec)

```
kubectl exec -it -n pg pg-postgresql-0 \
  -c postgresql -- /bin/sh
```

- Debug by copy

```
kubectl debug -it -n pg pg-postgresql-0 \
  --copy-to pg-postgresql-debug \
  --container debug-postgresql \
  --image busybox -- /bin/sh
```



Debug without Ephemeral Containers

- Debug by **attaching** limited by the tools in the image
- Debug by **copy** requires a pod restart
- Debugging workloads requires modifying number of replicas

Debug with Ephemeral Containers

- Debug by ephemeral container

```
kubectl debug -it -n pg pg-postgresql-0 \
  --image busybox -- /bin/sh
```



Debug with Ephemeral Containers

- Debug by ephemeral container

```
kubectl debug -it -n pg pg-postgresql-0 \
  --image busybox -- /bin/sh
```
- Allow sharing of container resources
- Ideal for complicated debugging of failing containers

Debug with Ephemeral Containers

```
$ kubectl debug \  
  -n pg \  
  -c debug-postgresql \  
  -it \  
  --image=busybox \  
  pod/pg-postgresql-0 \  
  -- /bin/sh
```

If you don't see a command prompt, try pressing enter.

```
# exit
```

Ephemeral Containers spec:

```
$ kubectl get pod -n pg pg-postgresql-0 -o yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2021-10-17T19:26:40Z"
...
spec:
  containers:
    ...
    ephemeralContainers:
    - image: busybox
      imagePullPolicy: Always
      name: debug-postgresql
    ...
status:
...
```



Ephemeral Containers spec:

...

spec:

containers:

...

ephemeralContainers:

- image: busybox
imagePullPolicy: Always
name: debug-postgresql

...

status:

...



Ephemeral Containers status:



```
$ kubectl get pod -n pg pg-postgresql-0 -o yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2021-10-17T19:26:40Z"
...
spec:
  ...
status:
  containerStatuses:
    ...
    ephemeralContainerStatuses:
    - containerID: docker://85a4eb6742bf ...
      image: busybox:latest
      imageID: docker-pullable://busybox@sha25 ...
  ...
```

Ephemeral Containers status:



...

spec:

...

status:

containerStatuses:

...

ephemeralContainerStatuses:

- containerID: docker://85a4eb6742bf ...
image: busybox:latest
imageID: docker-pullable://busybox@sha25 ...

...

What is a Container?

- What is a Container?



What is a Container?



- What is a Container?
- Containers can be thought of as Sandboxes
- Once in a Sandbox, cannot see what's outside of the Sandbox
- Controlling what a container can see is what isolates one Container from another



What is a Container?

- What is a Container?
- Containers can be thought of as Sandboxes
- Once in a Sandbox, cannot see what's outside of the Sandbox
- Controlling what a container can see is what isolates one Container from another
- **Linux Namespaces in the Node make this possible**

Linux Namespaces



- Linux Namespaces control Container isolation
- Linux Namespaces contain Container resources
 - processes, mounts, ...
- Linux Namespaces can be copied from Container to Container

Linux Namespaces

- There are many
 - mnt, pid, net, ipc, uts, uid, cgroup, time, ...
- Each namespace is identified by an inode



Linux Namespaces

- There are many
 - mnt, pid, net, ipc, uts, uid, cgroup, time, ...
- Each namespace is identified by an inode
- pid
 - Processes
- mnt
 - File system mounts
- net
 - Network interfaces





Linux Namespaces on node (docker)

```
$ docker ps # get docker container for node
```

CONTAINER ID	IMAGE	COMMAND ...	NAMES
abf011764870	kindest/node:v1.23.4	"/usr/local/bin/entr..."	kind-control-plane

```
$ docker exec -it kind-control-plane /bin/bash # exec into node
```

```
# ls -l /proc/1/ns/* # get namespaces for PID 1
```

```
lrwxrwxrwx 1 root root 0 May 7 20:12 /proc/1/ns/cgroup -> cgroup:[4026532459]
lrwxrwxrwx 1 root root 0 May 7 20:12 /proc/1/ns/ipc -> ipc:[4026532372]
lrwxrwxrwx 1 root root 0 May 7 20:12 /proc/1/ns/mnt -> mnt:[4026532370]
lrwxrwxrwx 1 root root 0 May 7 20:12 /proc/1/ns/net -> net:[4026532376]
lrwxrwxrwx 1 root root 0 May 7 20:00 /proc/1/ns/pid -> pid:[4026532374]
...
```

Linux Namespaces on node (docker)

```
# # get namespaces for PID 1
# ls -l /proc/1/ns/*
lrwx ... /proc/1/ns/cgroup -> cgroup:[4026532459]
lrwx ... /proc/1/ns/ipc -> ipc:[4026532372]
lrwx ... /proc/1/ns/mnt -> mnt:[4026532370]
lrwx ... /proc/1/ns/net -> net:[4026532376]
lrwx ... /proc/1/ns/pid -> pid:[4026532374]
```

Linux Namespaces on node (docker)

```
# pgrep -o postgres # get postgres PID
1935
```

```
# ps -f -p $(pgrep -o postgres) # get postgres process info
UID          PID      PPID  C  STIME          TIME CMD
1001         1935      1805  0  00:54      00:00:00 /opt/bitnami/postgresql ...
```

```
# ls -l /proc/$(pgrep -o postgres)/ns/* # list postgres namespaces
lrwxrwxrwx 1 1001 root 0 May  7 20:03 /proc/1946/ns/cgroup -> cgroup:[4026532921]
lrwxrwxrwx 1 1001 root 0 May  7 20:03 /proc/1946/ns/ipc -> ipc:[4026532917]
lrwxrwxrwx 1 1001 root 0 May  7 20:03 /proc/1946/ns/mnt -> mnt:[4026532919]
lrwxrwxrwx 1 1001 root 0 May  7 20:03 /proc/1946/ns/net -> net:[4026532837]
lrwxrwxrwx 1 1001 root 0 May  7 20:03 /proc/1946/ns/pid -> pid:[4026532920]
...
```

Linux Namespaces on node (docker)

```
# # list postgres namespaces
# ls -l /proc/$(pgrep -o postgres)/ns/*
lrwx ... /proc/1946/ns/cgroup -> cgroup:[4026532921]
lrwx ... /proc/1946/ns/ipc -> ipc:[4026532917]
lrwx ... /proc/1946/ns/mnt -> mnt:[4026532919]
lrwx ... /proc/1946/ns/net -> net:[4026532837]
lrwx .../proc/1946/ns/pid -> pid:[4026532920]
```


nsenter




- What you can see depends on what namespace you are in
- `nsenter` allows execution of a command within a namespace
- The command will only see what the namespace allows

nsenter

```
# nsenter --net=/proc/$(pgrep -o postgres)/ns/net \  
    ss --tcp -l
```

State	Recv-Q	Send-Q	Local Address:Port	Peer
LISTEN	0	224	0.0.0.0:postgresql	0.0.0.0:*
LISTEN	0	224	:::postgresql	:::*

nsenter



```
# nsenter --target $(pgrep -o postgres) --all ps -e
```

PID	TTY	TIME	CMD
1	?	00:00:01	postgres
137	?	00:00:00	postgres
138	?	00:00:00	postgres
139	?	00:00:00	postgres
140	?	00:00:00	postgres
141	?	00:00:00	postgres
142	?	00:00:00	postgres
22340	?	00:00:00	ps

Namespaces on Pod Container (K8s)

```
$ kubectl exec -it -n pg pg-postgresql-0 -- ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
1001	1	0	0	17:41	?	00:00:00	/opt/bitnami/postgresql ...
1001	137	1	0	17:41	?	00:00:00	postgres: checkpointer
1001	138	1	0	17:41	?	00:00:00	postgres: background writer
1001	139	1	0	17:41	?	00:00:00	postgres: walwriter
1001	140	1	0	17:41	?	00:00:00	postgres: autovacuum la ...
1001	141	1	0	17:41	?	00:00:00	postgres: stats collector
1001	142	1	0	17:41	?	00:00:00	postgres: logical repli ...
1001	1180	0	0	17:52	pts/0	00:00:00	ps -ef

Namespaces using nsenter

```
$ docker exec kind-control-plane /bin/bash -c -- \
  'nsenter --target $(pgrep -o postgres) --all /bin/ps -ef'
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
1001	1	0	0	17:41	?	00:00:00	/opt/bitnami/postgresql ...
1001	137	1	0	17:41	?	00:00:00	postgres: checkpointer
1001	138	1	0	17:41	?	00:00:00	postgres: background writer
1001	139	1	0	17:41	?	00:00:00	postgres: walwriter
1001	140	1	0	17:41	?	00:00:00	postgres: autovacuum la ...
1001	141	1	0	17:41	?	00:00:00	postgres: stats collector
1001	142	1	0	17:41	?	00:00:00	postgres: logical repli ...
1001	682	0	0	17:46	pts/1	00:00:00	/bin/sh
root	1073	0	0	17:50	pts/1	00:00:00	ps -ef

Ephemeral Containers: debug

- Network namespace shared by default



Ephemeral Containers

- Network namespace shared by default
- Other namespaces can be shared



Namespaces when debug by attaching

```
$ kubectl exec -it -n pg pg-postgresql-0 -c postgresql -- /bin/sh
```

```
# ls -l /proc/self/ns/*
```

```
1 1001 root 0 May 7 20:27 /proc/self/ns/cgroup -> cgroup:[4026532921]
1 1001 root 0 May 7 20:27 /proc/self/ns/ipc -> ipc:[4026532917]
1 1001 root 0 May 7 20:27 /proc/self/ns/mnt -> mnt:[4026532919]
1 1001 root 0 May 7 20:27 /proc/self/ns/net -> net:[4026532837]
1 1001 root 0 May 7 20:27 /proc/self/ns/pid -> pid:[4026532920]
1 1001 root 0 May 7 20:27 /proc/self/ns/pid_for_children -> pid:[4026532920]
1 1001 root 0 May 7 20:27 /proc/self/ns/time -> time:[4026531834]
1 1001 root 0 May 7 20:27 /proc/self/ns/time_for_children -> time:[4026531834]
1 1001 root 0 May 7 20:27 /proc/self/ns/user -> user:[4026531837]
1 1001 root 0 May 7 20:27 /proc/self/ns/uts -> uts:[4026532916]
```




KubeCon



CloudNativeCon

Europe 2022

Namespaces when using Ephemeral Container ...

```
$ kubectl debug -it -n pg pg-postgresql-0 --image busybox -- /bin/sh
```

Defaulting debug container name to debugger-n8bvb

If you don't see a command prompt, try pressing enter.

```
# ls -l /proc/self/ns/*
```

```
1 root root 0 May 7 20:28 /proc/self/ns/cgroup -> cgroup:[4026532924]
1 root root 0 May 7 20:28 /proc/self/ns/ipc -> ipc:[4026532917]
1 root root 0 May 7 20:28 /proc/self/ns/mnt -> mnt:[4026532922]
1 root root 0 May 7 20:28 /proc/self/ns/net -> net:[4026532837]
1 root root 0 May 7 20:28 /proc/self/ns/pid -> pid:[4026532923]
1 root root 0 May 7 20:28 /proc/self/ns/pid_for_children -> pid:[4026532923]
1 root root 0 May 7 20:28 /proc/self/ns/time -> time:[4026531834]
1 root root 0 May 7 20:28 /proc/self/ns/time_for_children -> time:[4026531834]
1 root root 0 May 7 20:28 /proc/self/ns/user -> user:[4026531837]
1 root root 0 May 7 20:28 /proc/self/ns/uts -> uts:[4026532916]
```

Share namespaces of target container ...



```
$ kubectl debug -it -n pg pg-postgresql-0 --image busybox \  
  --target postgresql -- /bin/sh
```

Defaulting debug container name to debugger-vfcxs

If you don't see a command prompt, try pressing enter.

```
# ls -l /proc/self/ns/*
```

```
1 root root 0 May 7 20:31 /proc/self/ns/cgroup -> cgroup:[4026532923]  
1 root root 0 May 7 20:31 /proc/self/ns/ipc -> ipc:[4026532917]  
1 root root 0 May 7 20:31 /proc/self/ns/mnt -> mnt:[4026532922]  
1 root root 0 May 7 20:31 /proc/self/ns/net -> net:[4026532837]  
1 root root 0 May 7 20:31 /proc/self/ns/pid -> pid:[4026532920]  
1 root root 0 May 7 20:31 /proc/self/ns/pid_for_children -> pid:[4026532920]  
1 root root 0 May 7 20:31 /proc/self/ns/time -> time:[4026531834]  
1 root root 0 May 7 20:31 /proc/self/ns/time_for_children -> time:[4026531834]  
1 root root 0 May 7 20:31 /proc/self/ns/user -> user:[4026531837]  
1 root root 0 May 7 20:31 /proc/self/ns/uts -> uts:[4026532916]
```

Ephemeral Container vs `kubectl exec`

- Share PID namespace
- Share Network Namespace
- Share IPC, UTS



Ephemeral Containers

- Ephemeral containers are a sub-resource





Sub-Resource

- Ephemeral containers are a sub-resource, "`ephemeralContainer`"
- Added to the pod spec `ephemeralContainers` property for each invocation
- Container exit status in the `ephemeralContainerStatuses` property
- Unique names must be used

Sub-Resource



```
$ kubectl get pod -n pg pg-postgresql-0 -o json | \
  jq '{"ephemeralContainers": [(.spec.ephemeralContainers[].name)], \
    "ephemeralContainerStatuses": \
    [(.status.ephemeralContainerStatuses[].name)]}'
{
  "ephemeralContainers": [
    "debug-postgresql",
    "debugger-n8bvb",
    "debugger-vfcxs"
  ],
  "ephemeralContainerStatuses": [
    "debug-postgresql",
    "debugger-n8bvb",
    "debugger-vfcxs"
  ]
}
```



Custom API calls

- Properties not available from `kubectl`
- Can mount container volumes
- Privileged containers
 - Set `securityContext`
- Cannot remove or change `ephemeralContainers` or `ephemeralContainerStatuses` property

Ephemeral Containers



```
$ kubectl proxy &
```

```
Starting to serve on 127.0.0.1:8001
```

```
$ curl -v -XPATCH -H "Content-Type: application/json-patch+json" \
  'http://127.0.0.1:8001/api/v1/namespaces/pg/pods/pg-postgresql-0/ephemeralcontainers' \
  --data-binary @- << EOF
```

```
[{
  "op": "add", "path": "/spec/ephemeralContainers/-",
  "value": {
    "command": [ "/bin/sh" ],
    "stdin": true, "tty": true,
    "image": "busybox",
    "name": "debug-container-1",
    "volumeMounts": [{
      "mountPath": "/mnt",
      "name": "data" }],
    "targetContainerName": "postgresql" }}}
```

```
EOF
```


Ephemeral Containers



```
$ kubectl proxy &
```

```
Starting to serve on 127.0.0.1:8001
```

```
$ curl -v -XPATCH -H "Content-Type: application/json-patch+json" \
  'http://127.0.0.1:8001/api/v1/namespaces/pg/pods/pg-postgresql-0/ephemeralcontainers' \
  --data-binary @- << EOF
```

```
[{
  "op": "add", "path": "/spec/ephemeralContainers/-",
  "value": {
    "command": [ "/bin/sh" ],
    "stdin": true, "tty": true,
    "image": "busybox",
    "name": "debug-container-1",
    "volumeMounts": [{
      "mountPath": "/mnt",
      "name": "data" }],
    "targetContainerName": "postgresql" }}}
```

```
EOF
```

Namespaces when debug by attaching

```
$ kubectl exec -it -n pg pg-postgresql-0 -c postgresql -- /bin/sh
```

```
# ls -l /proc/self/ns/*
```

```
1 1001 root 0 May 7 20:27 /proc/self/ns/cgroup -> cgroup:[4026532921]
1 1001 root 0 May 7 20:27 /proc/self/ns/ipc -> ipc:[4026532917]
1 1001 root 0 May 7 20:27 /proc/self/ns/mnt -> mnt:[4026532919]
1 1001 root 0 May 7 20:27 /proc/self/ns/net -> net:[4026532837]
1 1001 root 0 May 7 20:27 /proc/self/ns/pid -> pid:[4026532920]
1 1001 root 0 May 7 20:27 /proc/self/ns/pid_for_children -> pid:[4026532920]
1 1001 root 0 May 7 20:27 /proc/self/ns/time -> time:[4026531834]
1 1001 root 0 May 7 20:27 /proc/self/ns/time_for_children -> time:[4026531834]
1 1001 root 0 May 7 20:27 /proc/self/ns/user -> user:[4026531837]
1 1001 root 0 May 7 20:27 /proc/self/ns/uts -> uts:[4026532916]
```

Ephemeral Containers



```
$ kubectl exec -it -n pg pg-postgresql-0 -c debug-container-1 -- /bin/sh
```

```
# ls -l /proc/self/ns/*
```

```
1 root root 0 May 7 20:55 /proc/self/ns/cgroup -> cgroup:[4026532923]
1 root root 0 May 7 20:55 /proc/self/ns/ipc -> ipc:[4026532917]
1 root root 0 May 7 20:55 /proc/self/ns/mnt -> mnt:[4026532922]
1 root root 0 May 7 20:55 /proc/self/ns/net -> net:[4026532837]
1 root root 0 May 7 20:55 /proc/self/ns/pid -> pid:[4026532920]
1 root root 0 May 7 20:55 /proc/self/ns/pid_for_children -> pid:[4026532920]
1 root root 0 May 7 20:55 /proc/self/ns/time -> time:[4026531834]
1 root root 0 May 7 20:55 /proc/self/ns/time_for_children -> time:[4026531834]
1 root root 0 May 7 20:55 /proc/self/ns/user -> user:[4026531837]
1 root root 0 May 7 20:55 /proc/self/ns/uts -> uts:[4026532916]
```

Ephemeral Containers



```
# df # can mount pod volumes ("volumes" property in pod)
```

Filesystem	1K-blocks	Used	Available	Use%	Mounted on
overlay	123329088	61869768	55151536	53%	/
tmpfs	65536	0	65536	0%	/dev
/dev/vda1	123329088	61869768	55151536	53%	/mnt
/dev/vda1	123329088	61869768	55151536	53%	/etc/hosts
...					

```
# ls -l /mnt/data
```

```
PG_VERSION
base
global
...
postgresql.auto.conf
postmaster.opts
postmaster.pid
```

Netshoot



- Comprehensive toolset for diagnosis of network problems
- System level diagnosis tools
 - strace
 - ltrace
 - tcpdump
 - ...

Ephemeral Containers



```
$ curl -v -XPATCH -H "Content-Type: application/json-patch+json" \
  'http://127.0.0.1:8001/api/v1/namespaces/pg/pods/pg-postgresql-0/ephemeralcontainers' \
  --data-binary @- << EOF
[{"op": "add", "path": "/spec/ephemeralContainers/-",
  "value": {
    "command": [ "/bin/sh" ],
    "stdin": true, "tty": true,
    "image": "nicolaka/netshoot",
    "name": "debug-container-2",
    "securityContext": { "privileged": true },
    "volumeMounts": [{
      "mountPath": "/mnt",
      "name": "data" }],
    "targetContainerName": "postgresql" }]}
EOF
```

Ephemeral Containers



```
$ curl -v -XPATCH -H "Content-Type: application/json-patch+json" \
  'http://127.0.0.1:8001/api/v1/namespaces/pg/pods/pg-postgresql-0/ephemeralcontainers' \
  --data-binary @- << EOF
[{
  "op": "add", "path": "/spec/ephemeralContainers/-",
  "value": {
    "command": [ "/bin/sh" ],
    "stdin": true, "tty": true,
    "image": "nicolaka/netshoot",
    "name": "debug-container-2",
    "securityContext": { "privileged": true },
    "volumeMounts": [{
      "mountPath": "/mnt",
      "name": "data" }],
    "targetContainerName": "postgresql" }]}
EOF
```

Ephemeral Containers

```
# strace -p $(pgrep postgresql) 2>&1 | head
```

```
strace: Process 1 attached
```

```
select(8, [5 6 7], NULL, NULL, {tv_sec=55, tv_usec=50862 ...
```

```
rt_sigprocmask(SIG_SETMASK, ~[ILL TRAP ABRT BUS FPE SEGV ...
```

```
accept(5, {sa_family=AF_INET, sin_port=htons(36538), sin_...
```

```
getsockname(9, {sa_family=AF_INET, sin_port=htons(5432), ...
```

```
# exit
```




KubeCon



Google Cloud

Namespaces when debug by attaching (slide 40)

```
$ kubectl exec -it -n pg pg-postgresql-0 -c postgresql -- /bin/sh
```

```
# ls -l /proc/self/ns/*
```

```
1 1001 root 0 May 7 20:27 /proc/self/ns/cgroup -> cgroup:[4026532921]
1 1001 root 0 May 7 20:27 /proc/self/ns/ipc -> ipc:[4026532917]
1 1001 root 0 May 7 20:27 /proc/self/ns/mnt -> mnt:[4026532919]
1 1001 root 0 May 7 20:27 /proc/self/ns/net -> net:[4026532837]
1 1001 root 0 May 7 20:27 /proc/self/ns/pid -> pid:[4026532920]
1 1001 root 0 May 7 20:27 /proc/self/ns/pid_for_children -> pid:[4026532920]
1 1001 root 0 May 7 20:27 /proc/self/ns/time -> time:[4026531834]
1 1001 root 0 May 7 20:27 /proc/self/ns/time_for_children -> time:[4026531834]
1 1001 root 0 May 7 20:27 /proc/self/ns/user -> user:[4026531837]
1 1001 root 0 May 7 20:27 /proc/self/ns/uts -> uts:[4026532916]
```

Ephemeral Containers



```
$ kubectl exec -it -n pg pg-postgresql-0 -c debug-container-2 -- /bin/sh
```

```
# ls -l /proc/self/ns/*
```

```
1 root root 0 May 7 21:01 /proc/self/ns/cgroup -> cgroup:[4026532459]
1 root root 0 May 7 21:01 /proc/self/ns/ipc -> ipc:[4026532917]
1 root root 0 May 7 21:01 /proc/self/ns/mnt -> mnt:[4026532924]
1 root root 0 May 7 21:01 /proc/self/ns/net -> net:[4026532837]
1 root root 0 May 7 21:01 /proc/self/ns/pid -> pid:[4026532920]
1 root root 0 May 7 21:01 /proc/self/ns/pid_for_children -> pid:[4026532920]
1 root root 0 May 7 21:01 /proc/self/ns/time -> time:[4026531834]
1 root root 0 May 7 21:01 /proc/self/ns/time_for_children -> time:[4026531834]
1 root root 0 May 7 21:01 /proc/self/ns/user -> user:[4026531837]
1 root root 0 May 7 21:01 /proc/self/ns/uts -> uts:[4026532916]
```



Linux Namespaces on node (docker) (slide 30)

```
# # get namespaces for PID 1
```

```
# ls -l /proc/1/ns/*
```

```
lrwxr ... /proc/1/ns/cgroup -> cgroup:[4026532459]
```

```
lrwxr ... /proc/1/ns/ipc -> ipc:[4026532372]
```

```
lrwxr ... /proc/1/ns/mnt -> mnt:[4026532370]
```

```
lrwxr ... /proc/1/ns/net -> net:[4026532376]
```

```
lrwxr ... /proc/1/ns/pid -> pid:[4026532374]
```

Distroless Containers



- Separate production release from tools
- Prod and tools projects become version independent
 - Can select tools appropriate to task
 - No longer dependent on versions in container
- Distroless containers are faster to build
 - Fewer layers

Notes



<https://opensource.googleblog.com/2022/01/Introducing+Ephemeral+Containers.html>

<https://www.ianlewis.org/en/what-are-kubernetes-pods-anyway>

<https://icicimov.github.io/blog/virtualization/Linux-Container-Basics/>

<https://bit.ly/38Nvd6P>

https://en.wikipedia.org/wiki/Linux_namespaces

<https://en.wikipedia.org/wiki/Cgroups>

<https://github.com/GoogleContainerTools/distroless>

<https://kind.sigs.k8s.io>

<https://github.com/nicolaka/netshoot>

<https://github.com/eldadru/ksniff>